



CURIO *a starter*

volume *i*
five experiments



the curio starter

Five things to try with Claude this week.

*Copy-paste prompts.
No theory.
Just things to try.*

*From the upcoming book Ask the Mirror — The
Coming Age of the AI Oracle.
Ryan O'Connor, Curio.*

first edition
curio.help





Five things to try with Claude this week.

One a day, Monday through Friday. Ninety minutes total.

Curio is a guide for the AI uncertain — for the senior creative operators who can feel the wave but can't quite see their part in it. The site at curio.help walks the long version. This doc walks the short version: **five concrete experiments**, plus a five-minute Sunday-night setup that makes all of them land harder.

Claude is what we recommend (and what most of this is written for). ChatGPT works. Gemini works. By Friday afternoon you'll know which of these moves are about to change your work, and which ones aren't for you yet. Both are useful answers.

how to use this doc

Open it on your second monitor. **Sunday night:** the setup (page 4) and the relationship skill (page 5). Ten minutes, tops. **Monday through Friday:** one experiment a day. The shittier and more specific your first draft, the better the answer comes back.

from the book this comes from

Ask the Mirror — The Coming Age of the AI Oracle. The long argument for why these moves work. Publishing later this year via Curio. Waitlist on the back page.

— Ryan O'Connor





Five reasons people stall on AI.

If you took the Curio quiz, you already know yours. If you didn't, here they are in one breath each.

01 Exhaustion

"I don't have time."

*The tool **is** the time.*

02 Guilt

"Is this cheating?"

The room already moved.

03 Identity

"I'm the Excel person."

*You are **more** than your tool.*

04 Muscle Memory

"Someday I'll be good at AI."

There is no someday.

05 Yesterday

"I tried it eight months ago and it was wrong."

It compounds. You don't — unless you keep asking.

Each experiment in this doc breaks one of them.





Before you start.

Five minutes. Sunday night. So Monday's clean.

01

Open an AI.

Claude is what we recommend — claude.ai. ChatGPT works — chatgpt.com. Gemini works — gemini.google.com. Pick one. Sign in. Open it in a fresh tab.

02

Spend the twenty bucks.

The free tier is fine for a taste. For the real moves you want a paid model — *Claude Pro* or *ChatGPT Plus*, twenty dollars a month. It's the cheapest senior colleague you'll ever hire.

03

Start a project.

Claude has Projects. ChatGPT has Projects. Make one called *Curio Starter* — or just your name. Every experiment this week lives in that same project so context accumulates instead of evaporating.

*Then turn the page and run the relationship skill — **once**, before Monday opens.*

five minutes — tops





Build the relationship.

Skip this and the rest works at half strength.

the setup

Most people open Claude (or ChatGPT) and start asking for things. That's like meeting a senior colleague on day one and immediately handing them a task with no context. They can do it. But they'll do it at maybe thirty percent of what they're actually capable of.

Spend ten minutes building the relationship. Run the prompt on the next page. Answer the questions like you're meeting a smart colleague for coffee. By the end you'll have two short files that you re-paste into every conversation after.

why this matters

Every experiment in this doc lands harder when the AI already knows you. The pantry move is sharper. The "i wish" move is more accurate. The hate inventory targets the right tasks. The relationship is the move under all the other moves.

two ways to run this

Option A — the prompt. Turn the page, copy the prompt, paste it in. Works in Claude, ChatGPT, Gemini, anywhere. Takes ten minutes. The full interview, the email calibration, and your two files at the end.

Option B — the skill. If you're on Claude, we packaged this whole interview as a runnable skill. Download *curio-relationship.skill* from *curio.help/skill*, drop it into your Project's Skills folder, then say "*run the curio-relationship skill*". Same interview. Same outputs. One-tap re-run later when your work shifts.

The move under all the others · teach it who you are first.

once is enough





↳ the relationship prompt – copy this whole block

interview me – like a smart new colleague would
on day one. friendly, curious, one question at
a time. no jargon. follow the thread when i give
you something interesting.

learn enough about me to be useful.

at the end, give me two short files i can save
and re-paste into every future conversation:

1. **About-me.md** – who i am, what i do, what i'm
working on right now, what i'd like to be
doing more of, who i work with day to day.
2. **AI-style.md** – how i want you to talk to me,
whether to push back or just deliver, what i
don't want (fluff, apologies, jargon, hedging),
the rhythm we keep across our conversations.

ask your first question.

what to do with the result

When the interview wraps, copy *About-me.md* and *AI-style.md*
somewhere you can find them, then drop them into the Project
you made on page four. From here on, every new conversation
opens with an AI that already knows you. Monday lands harder
because of this.

Claude users: the packaged version is at curio.help/skill. Same
interview, runnable on demand. Includes a short email-writing
exercise that calibrates how you actually write.

Try this with: a Sunday night and a cup of coffee.

ten minutes – once





01

Test the bias.

Run yesterday's question, today.

the setup

There's a question you asked AI six or more months ago. The answer disappointed you. You haven't asked it since. That's a bias — possibly the most expensive one in your career right now. Today's model isn't yesterday's. Today's answer probably isn't either.

```
↳ the prompt – copy this whole block
```

```
six months ago i tried this question with AI:
```

```
[PASTE THE QUESTION YOU ASKED THEN]
```

```
back then it gave me [DESCRIBE WHAT WAS WRONG  
WITH THE ANSWER].
```

```
today i want to:
```

1. ask the exact same question again – answer it now, in full
2. tell me how the answer differs from a typical model response in early/mid 2024
3. tell me what's improved most: reasoning, accuracy, speed, or something else

```
[ASK THE QUESTION HERE – full text]
```

what to do with the result

Notice the difference. Notice you just updated your prior in thirty seconds. Notice you might have been wrong for six months — and that this is the cheapest possible way to find out.





02

Stop asking how. Hand Claude the pantry.

the setup

You're working on a project. Normally you'd start by asking "*help me write a [deck / email / strategy / brief].*" That's asking the AI to be a tool.

Instead, hand it the *pantry* — every doc, every note, every scrap of context you have — and ask it to make you the dish.

why this works

Most people micromanage their AI. They scope, they outline, they ask permission. The pantry move treats it like a senior colleague — give it everything, let it judge what matters, react to its draft.

You go from blank page to seventy-percent-there in four minutes. That's the whole game.

The fallacy this breaks · Muscle Memory.

there is no someday





```
↳ the prompt – copy this whole block
```

```
i'm working on [PROJECT NAME] for [CLIENT / CONTEXT].
```

```
here are all the ingredients i have:
```

```
[paste in everything – briefs, decks, notes,  
emails, slack threads, research, prior work,  
anything tangentially related. don't filter.]
```

```
i need a [DELIVERABLE – e.g., "client-facing  
proposal," "internal strategy memo," "one-pager  
for the team"].
```

```
don't ask me clarifying questions.
```

```
make me a first draft.
```

```
i'll tell you what's wrong.
```

what to do with the result

Read it. Notice what's right. Edit what's wrong. The point isn't that the draft is perfect — it's that you went from blank-page to seventy-percent-there in the time it took to make a coffee.

Then react. Tell it what was off. Watch the second draft. That's where the move pays for itself.

Try this with: the next deliverable on your list.

the big hookup





03

List three things you hate doing. Hand them over.

the setup

There are tasks you do every week that drain you. Not because they're hard — because they're boring, repetitive, soul-corroding.

You've been doing them by hand because *"it'd take longer to teach AI to do it than to just do it."* That math has changed.

why this works

AI shines on the boring middle of your week. Most people apply it to the glamorous tasks — writing, strategy, big-think. The bigger return is killing the things you don't even notice draining you.

One to three hours per task per week is normal. Stack three tasks and you've bought yourself a Friday afternoon.

The fallacy this breaks · Exhaustion.

the tool is the time





```
↳ the prompt – copy this whole block
```

here are three things i hate doing every week:

1. [SPECIFIC TASK – e.g., "writing client status update emails on Fridays"]
2. [SPECIFIC TASK – e.g., "summarizing meeting notes into action items"]
3. [SPECIFIC TASK – e.g., "formatting our weekly team digest"]

for each one:

- show me a workflow where you do 80% of it
- write a saved prompt or skill i can run weekly
- estimate the time it'd save me per week

start with #1.

what to do with the result

Pick the one with the highest time savings. Use the saved prompt next week. Time it.

Most people save one to three hours per task per week. Do this for three tasks and you've made yourself an afternoon back, every week, forever.

Try this with: the chore that keeps showing up on your Friday list.

compounds weekly





04

Don't redo the work. Teach the AI to do the work.

the setup

Last Tuesday you spent thirty minutes getting your AI to help with something. Next Tuesday you'll do the same thing — same prompt, same effort. **Stop.** Run this right after you finish any task you'll repeat.

```
4 run after you've used your ai to do a task once
```

```
the task we just did — i'm going to do this again.  
probably weekly.
```

```
based on what we just worked through, write me:
```

1. a "saved prompt" version i can paste in next week
to skip the setup
2. a "skill" version that wraps this whole workflow
into one command
3. a "trigger guide" — when should i actually run this?
what's the signal that says "time to run the saved prompt"?

```
i want next week's version of this task to take  
2 minutes, not 30.
```

why this is the move

Most people use their AI like a chat — repeat tomorrow, same effort. The fluent ones save the work, not the chat. Time savings stack like interest.

Run this once a week. In ninety days you'll have a private library of saved prompts that take two minutes each. Compounding is the whole game.





05

Stop asking how do I. Start with I wish.

the setup

There's something you've been trying to do for weeks. You don't know how. You've been Googling. You've been asking peers. You haven't asked your AI — because you'd have to figure out the right "*how do I*" first.

Skip that step.

why this works

AI is built on language. "*How do I*" assumes you already know the shape. "*I wish*" assumes you only know the desire. The second one is closer to the truth.

By Friday, four days in, your AI knows you. It can take a wish further than it could on Monday. End the week here.

*The fallacy this breaks · **Mindset.***

stop figuring out how





```
↳ the prompt - copy this whole block
```

```
i wish i had a tool that did [SPECIFIC THING].
```

```
i don't know how to build it.
```

```
i don't know what to call it.
```

```
i don't know what stack to use.
```

```
but i know what it should do:
```

```
[describe in plain english - 3 to 5 sentences.
```

```
who would use it. what problem it solves.
```

```
what the moment of using it looks like.]
```

```
build me a working prototype.
```

```
don't ask permission.
```

```
show me what it would look like.
```

what to do with the result

Sometimes you get a working app. Sometimes you get a plan and a first version of code. Sometimes you get a doc walking you through it. Whatever you get — try it. Iterate.

The thing that matters isn't whether the first version works. It's that you stopped guarding the desire and let it become a thing in the world. That's the end of the week.

Try this with: that idea you've been carrying around for three months.

no permission required





You ran the experiments. What now?

If any of those hit, the rest of Curio is built to take you further.

free · weekly

The Substack.

One essay every Tuesday. The shape and voice of this doc, extended. No paywall. No spam. curioco.substack.com

free · the
long version

The four focuses.

The site walks the long version of these ideas through four focuses — *Inspiration*, *Filter*, *Mindset*, *Solutions* — each with its own reframe, examples, and deeper prompts. curio.help

\$9 · the new
user

\$9 *The First Hour with Claude.*

A full setup walkthrough and six more starter skills. For the operator who just opened Claude and wants to skip the first month of figuring it out.

ryanoconnor5.gumroad.com/l/curio_firsthour

\$25 · the
year-in user

\$25 *The Field Guide.*

The longer playbook for the operator already using AI a year or more and quietly stopped getting better at it. ryanoconnor5.gumroad.com/l/curio_fieldguide

on the way

Ask the Mirror — the book.

The long argument behind everything you just read. Publishing later this year.

curio.help/book · join the waitlist.

\$200–\$1,000

A real conversation.

Bounded consult, scoped to one question. Three turns of email + one Zoom + a tangible output. Inbound only — reply to any Substack post and tell me what you're trying to figure out.

*Stop asking for help.
F'ing do it.*

— ryan o'connor · curio

