

The Contract You Didn't Know You Signed

I was onboarding a new team member last week. Sharp. Experienced. The kind of person who runs a huge client account on instinct and institutional memory, stays late, carries more than her share. When I brought up AI, she said the thing I've heard a dozen times now:

"I don't have time to learn a new thing."

I've learned not to argue with that sentence. Because she's right — if Claude were a new thing to learn, she genuinely wouldn't have time. But it isn't. And the fastest way to show her that is to show her something that's already hers.

So I pulled up a dashboard I'd built. Live data. Her kind of layout. She leaned in.

Then something interesting happened. She started thinking out loud about how she could split two program tracks, pull from Asana, build views for different stakeholders. She was already scoping the project. Already breaking it into tasks. Already figuring out how to execute it herself.

I stopped her.

You don't have to do any of that. You just tell it what you want.

She looked at me like I'd said something slightly absurd. And I realized — that reaction isn't resistance. It's conditioning. Thirty years of it.

Here's what I've been thinking about since.

For three decades, the relationship between humans and computers has had one fundamental shape: **you learn the tool's language, the tool executes your commands**. Every advancement in that time — faster processors, better interfaces, cloud integration, mobile — was an improvement within that shape. The human still adapts. The machine still executes. That's the deal.

Even the most sophisticated software of the last decade works this way. Salesforce, Notion, Excel, project management platforms — all of them require you to speak their language. You find the feature. You build the formula. You figure out the workaround. You accept what the tool can and can't do and you adjust your expectations accordingly.

That adjustment — that quiet, constant negotiation between what you actually want and what the software will allow — is so habitual most people don't notice they're doing it. It just feels like working.

She told me she'd tried Notion. Tried colored dots in Outlook. "They're one-way," she said. She understood, correctly, that every tool she'd ever used was asymmetric. She gives. It does its fixed thing. She adjusts. That's the deal.

I'd call this a cognitive contract. And almost everyone in knowledge work signed it sometime in the early 90s and has been honoring it ever since.

The contract reads: *I will scope my needs to what the tool can do. I will adapt my workflow to the tool's logic. I will accept the tradeoffs as the price of using technology.*

The last time this contract was seriously renegotiated was the leap from analog to digital. But even that leap was less radical than it looked. Writing a memo on a typewriter versus writing one in Word — same thinking, faster execution, easier editing. The computer amplified output. It didn't change the nature of the relationship. You were still the one doing the cognitive work. The computer did exactly what you told it, nothing more, nothing less.

What's different now isn't speed. It isn't features. It's agency.

For the first time, the tool can understand intent, not just instruction. You don't have to translate your need into the software's logic. You say what you want — vaguely, incompletely, in your own words — and it figures out the rest.

That is not an upgrade within the same category. That's a different category entirely.

And the psychological problem is that thirty years of conditioning makes people reach for the wrong mental model automatically. They ask: *What does this software do? What's the learning curve? Where's the manual? What are the limitations?* Those are exactly the right questions for every tool that came before. They are the wrong questions for this.

The right questions are: *What do I actually want? What's been draining me? What would my work look like if I didn't have to make any tradeoffs?*

Most people haven't asked themselves those questions about work in years. They forgot what they actually wanted — because wanting something a tool couldn't do was pointless. You learned not to want it.

She asked me: *But what would I even wish for? I don't know what I'd ask.*

That's not a blocker. That's actually the best place to start.

Start with pain, not wishes. Most people can name what's draining them far faster than they can articulate an ideal state. Describe the problem. Let it propose the solution. React, refine, redirect. You're curating — not building from scratch.

I told her: think about what you want your dashboard to look like. Your style. Your data. The way your brain works, not the way the software works. Bullet points, summaries, color codes — whatever.

She laughed. *Anime characters?*

Sure, I said. Whatever you want.

That was the moment. Not a demo. Not a tutorial. Just the realization that the negotiation she'd been doing her whole career — what she wanted versus what was possible — was over. The tool bends now. To her. Not the other way around.

The leap from analog to digital changed how fast we worked. This changes who does the work.

We are moving from a world where humans learn the machine's language — to a world where the machine learns ours. That shift is harder to see than a new interface or a faster processor. It lives in the moment you stop scoping and start directing. When you catch yourself about to break something into tasks and realize you don't have to. When you ask for what you actually want instead of what you think is achievable.

The cognitive contract is expiring. Most people haven't noticed yet.

The ones who do — who make that shift, even partially, even messily — are going to look back at how everyone else is still working and feel exactly what that person from 1988 would feel watching us today.

Disbelief. Then wonder.

Ryan O'Connor is VP of Creative & Innovation at RedPeg Marketing, where he leads the agency's AI adoption and creative operations.